Introduction to Machine Learning and Python

Shrey Gupta

Applied Machine Learning (HOUSECS 59-01), Duke University

August 29, 2018

Topics

- **Classification:** naive Bayes, support vector machines, kernel methods, and neural networks.
- **Regression:** spline interpolation and linear and polynomial regression.
- **Unsupervised learning:** mixture of Gaussians clustering.
- **Computer vision:** object detection via convolutional neural networks, feature extraction, edge detection, and processing methods.
- **Dimensionality reduction:** generalized discriminant analysis.
- **Evaluation of machine learning models.**

Prerequisites

- This is an *applied* course, but requires fundamental understanding of the algorithms and techniques being used.
- Prerequisites: basic fluency in programming and mathematics at the single-variable calculus level.

Grading

- **Attendance:** required to attend at least 11 (of 14) classes. 11 classes are lectures, and three are office hours.
- **Programming assignments:** required to complete all three assignments, with individual scores of 70% or greater.
- **Final paper:** 1500-word write-up detailing machine learning model from final programming assignment.
- Class is graded on a *satisfactory/unsatisfactory* basis.
- See syllabus for more detailed information.

What is Machine Learning?

- "Give computers the ability to learn without being explicitly programmed" (Arthur Samuel).
- Formal problem specification: "A computer program is said to learn from experience E with respect to some class of tasks T and performance measure P if its performance at tasks in T, as measured by P, improves with experience E" (Tom Mitchell).

Terminology

- Grouped into two categories: *supervised* and *unsupervised* learning.
- Input data comes in *features* that describe each data point.
- *Training data*: data used to train algorithm (i.e. create model).
- *Testing data*: data used to evaluate performance of algorithm.

Supervised Learning

- Data (a subset from a larger distribution) is labeled, and we attempt to generalize to (predict) the larger distribution.
- Regression: predicts a continuous value output.
- Classification: predicts a discrete class output.

Regression: Examples

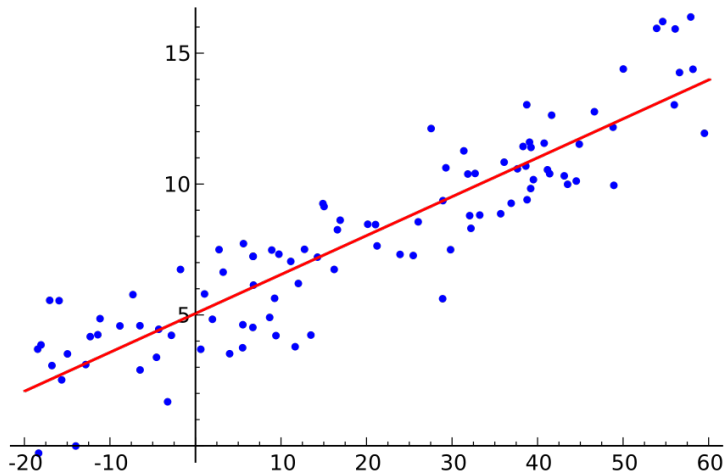

- Given data about square footage, age, zip code, and housing demand, predict the selling price for a house.
- Predict the percentage increase or decrease in the price of an equity.

Classification: Examples

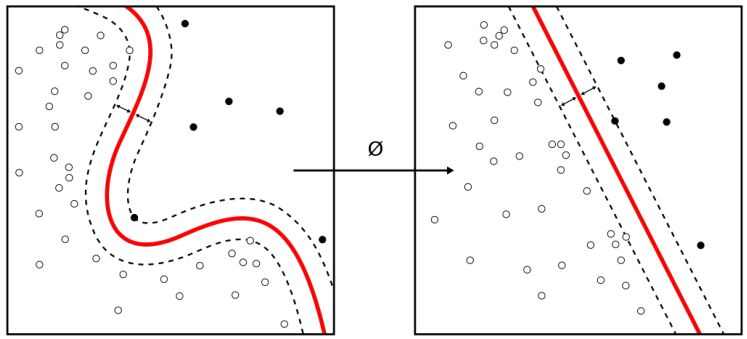

- Given data about temperature, humidity, and wind speed, predict whether it will be sunny, cloudy, or raining.
- Predict whether the price of an equity will increase or decrease.

Image source: Wikipedia

Unsupervised Learning

- Data is unlabeled (no "ground truth").
- Problems: clustering, density estimation, and pattern detection.

Clustering: Examples

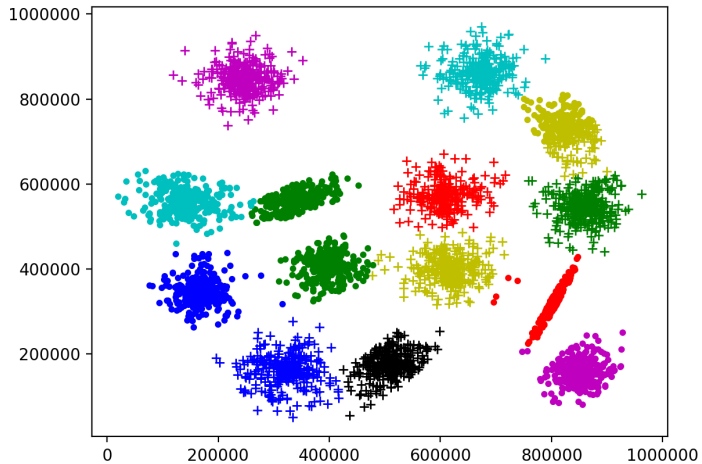

- Given consumption data, partition the consumers into market segments.
- Given several news articles (and their text), group them based on similarity.

Python

- We'll be using *Python 3.x* throughout the course.
- Libraries and frameworks: NumPy, SciPy, Pandas, Matplotlib, SciKit, TensorFlow, NLTK, and Jupyter.
- Today's notebook will ensure all of the packages are downloaded and work through an introduction of Python.