

# Naive Bayes Classifiers

Shrey Gupta


*Applied Machine Learning* (HOUSECS 59-01), Duke University

October 3, 2018

# Supervised Learning

- ▶ Data (a subset from a larger distribution) is labeled, and we attempt to generalize to (predict) the larger distribution.
- ▶ Classification: predicts a discrete class output.

# Classification: Examples



- ▶ Given data about temperature, humidity, and wind speed, predict whether it will be sunny, cloudy, or raining.
- ▶ Predict whether the price of an equity will increase or decrease.

*Image source: Wikipedia*

# Recall

$$X = \begin{bmatrix} x_1^{(1)} & x_1^{(2)} & x_1^{(3)} & \dots & x_1^{(m)} \\ x_2^{(1)} & x_2^{(2)} & x_2^{(3)} & \dots & x_2^{(m)} \\ \vdots & & & \ddots & \vdots \\ x_n^{(1)} & x_n^{(2)} & x_n^{(3)} & \dots & x_n^{(m)} \end{bmatrix}, y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$$

- ▶ Data is stored in matrices and vectors.
- ▶ Given  $n$  (training) data points and  $m$  features (per data point).
- ▶ Given labeled data vector  $y$ .

## Recall

$$X_{test} = \begin{bmatrix} x_1^{(1)} & x_1^{(2)} & x_1^{(3)} & \dots & x_1^{(m)} \\ x_2^{(1)} & x_2^{(2)} & x_2^{(3)} & \dots & x_2^{(m)} \\ \vdots & & & \ddots & \vdots \\ x_k^{(1)} & x_k^{(2)} & x_k^{(3)} & \dots & x_k^{(m)} \end{bmatrix}, \hat{y}_{test} = \begin{bmatrix} \hat{y}_1 \\ \hat{y}_2 \\ \vdots \\ \hat{y}_k \end{bmatrix}$$

- ▶ Given  $k$  testing data points and  $m$  features (per data point).
- ▶  $\hat{y}_{test} = f(X_{test})$  contains *predictions* of the classification algorithm, where  $f(\cdot)$  is learned by the algorithm.
- ▶ How do we define  $f(\cdot)$ , and how does the algorithm “learn” it?

# Introduction

- ▶ The Naive Bayes algorithm is a simple probabilistic classification algorithm.
- ▶ It is “naive” in the sense that it assumes features to be *independent* and *equal*. Note this doesn't always hold true, but the algorithm often works well in practice.
- ▶ The argument is largely based on *Bayes' theorem*, a crucial probability theorem.

# Notation

- ▶ Define  $P(A)$  to be the *probability* of an event  $A$ .
- ▶ Let  $P(A|B)$  be the *conditional probability* of event  $A$  on event  $B$ . That is, assuming event  $B$  happens, what is the probability of event  $A$ ?

# Goal

$$P(y|X) = P(y|x^{(1)}, x^{(2)}, \dots, x^{(m)})$$

$$\hat{y} = f(X) = \arg \max_y P(y|X)$$

- ▶ What is the probability of class  $y$  given features  $x^{(1)}, x^{(2)}, \dots, x^{(m)}$ ?
- ▶ Our prediction  $\hat{y}$  will be the class  $y$  that maximizes such probability.




## Optional: Bayes' Theorem

$$P(A|B) = \frac{P(B|A) \cdot P(A)}{P(B)} = \frac{\textit{likelihood} \times \textit{prior}}{\textit{evidence}}$$

- ▶ Prior:  $P(A)$  is the probability of event  $A$  prior to observing any “evidence”.
- ▶ Likelihood:  $P(B|A)$  is the likelihood of observing evidence  $B$  given event  $A$ .
- ▶ Evidence:  $B$  is the evidence observed.
- ▶ Posterior:  $P(A|B)$  is the probability of event  $A$  after observing evidence  $B$ .

## Optional: Theorem Visualization



*Image source: Oscar Bonilla*

## Optional: Theorem Derivation

$$\begin{aligned} P(A|B) &= \frac{P(A|B) \cdot P(B)}{P(B)} \\ &= \frac{P(A \cup B)}{P(B)} = \frac{P(B|A) \cdot P(A)}{P(B)} \end{aligned}$$

## Optional: Algorithm Derivation

$$\begin{aligned} P(y|X) &= \frac{P(X|y) \cdot P(y)}{P(X)} \\ &\propto P(X|y) \cdot P(y) \end{aligned}$$

- ▶  $P(X)$  can be treated as a constant across all different  $y$  under examination, and hence can be “discarded”.
- ▶ Therefore, our goal is to compute  $P(X|y) \cdot P(y) \forall y$ .

## Optional: Algorithm Derivation

$$\begin{aligned} P(X|y) \cdot P(y) &= P(X, y) = P(x^{(1)}, x^{(2)}, \dots, x^{(m)}, y) \\ &= P(x^{(1)}|x^{(2)}, \dots, x^{(m)}, y) \cdot P(x^{(2)}, \dots, x^{(m)}, y) \\ &= P(x^{(1)}|x^{(2)}, \dots, x^{(m)}, y) \cdot P(x^{(2)}|x^{(3)}, \dots, x^{(m)}, y) \cdot P(x^{(3)}, \dots, x^{(m)}, y) \\ &= P(x^{(1)}|x^{(2)}, \dots, x^{(m)}, y) \cdots P(x^{(m)}|y) \cdot P(y) \end{aligned}$$

- ▶ We can “naively” assume conditional independence between features so that  $P(x^{(i)}|x^{(i+1)}, \dots, x^{(m)}, y) = P(x^{(i)}|y)$ .

## Optional: Algorithm Derivation

$$\begin{aligned} & P(x^{(1)}|x^{(2)}, \dots, x^{(m)}, y) \cdots P(x^{(m)}|y) \cdot P(y) \\ &= P(x^{(1)}|y) \cdot P(x^{(2)}|y) \cdots P(x^{(m)}|y) \cdot P(y) \\ &= P(y) \cdot \prod_i P(x^{(i)}|y) \end{aligned}$$

- We can normalize over all classes  $y_j$  to find the probability  $P(y|X)$ . More formally,

$$P(y|X) = \frac{P(y) \cdot \prod_i P(x^{(i)}|y)}{\sum_j P(y_j) \cdot \prod_i P(x^{(i)}|y_j)}.$$

# Practicalities

- ▶ In practice, it is not necessary to normalize and compute  $P(y|X)$ .
- ▶ Remember:  $P(X)$  is effectively a constant across all  $y$ , so we only need to maximize the numerator.

# Algorithm

$$g(X, y) = P(y) \cdot \prod_i P(x^{(i)}|y)$$

$$\hat{y} = f(X) = \arg \max_y g(X, y)$$

- ▶ Compute *score*  $g(X, y)$  over classes  $y$ .
- ▶ Assign prediction  $\hat{y}$  to the class  $y$  that maximizes such score.
- ▶  $P(y)$  and  $P(x^{(i)}|y)$  can be easily calculated using the data.  
However, this assumes the features are *discrete*.



# Continuous Features

$$P(x^{(i)}|y) = \frac{1}{\sqrt{2\pi(\sigma_y^{(i)})^2}} \cdot e^{-\frac{(x^{(i)} - \mu_y^{(i)})^2}{2(\sigma_y^{(i)})^2}}$$

- ▶ If a feature is instead *continuous* and we assume it to be *normally distributed*, we can calculate the feature's mean  $\mu_y^{(i)}$  and variance  $(\sigma_y^{(i)})^2$  over class  $y$ .
- ▶  $P(x^{(i)}|y)$  can then be calculated using the above formula.

# Notebook

- ▶ We've introduced several theoretical concepts today, so some (or a lot) of it might be confusing.
- ▶ Today's notebook will work through an example of the Naive Bayes classifier with discrete and continuous features.